

ComicEngine: A Code-First Pipeline for Consistent, Cost-Aware AI History Comics

Technical Report on the *onceuponatime* System

Living Phase Log (updated after each phase gate)

Snapshot: 27 August 2026 · Phases 0–11 + Romance HIMYM Ep1 + Version 2B B0–B6 + G1

Avinash Nandyala

Independent Project / ComicMainEngine

github.com/navinash47/ComicMainEngine

Report regenerate: `python scripts/build_report.py`

Abstract

We present **ComicEngine**, a modular, code-first system for generating educational history comics with recurring characters, tracked API spend, and human-in-the-loop curation. Built for *onceuponatime* (bedtime narratives that explain real events to children), the system separates *coding/LLM* traffic (optionally via OmniRoute at `localhost:20128`) from *image generation* (direct commercial APIs: fal FLUX, Google Gemini image). This living report records what each phase actually ran, what worked, quantitative ledger totals from SQLite (`data/usage.db`), style bake-off judge scores, and interim character-consistency findings. **As of this snapshot:** Phase 9 adds an ROI dashboard at `/roi` with unit economics (\$/story, \$/panel, 100-episode projection), spend breakdowns by phase/provider/category/purpose, charts, and model recommendations from the live ledger ($\approx$ \$2.96 R&D so far).

1 Introduction

Production AI comics need character/style consistency, predictable cost, commercially licensed image paths, and an auditable spend ledger. ComicEngine addresses these with phasal CLIs, Pydantic episode JSON, style presets, and a FastAPI dashboard (Tasks, Stories, Analytics, Images, and this Report PDF at `/report`).

Test story. End-to-end stress test: dramatized bedtime episode about the Cockroach Janta Party (CJP) / fair-exam student protests (characters including Abhijeet Dipke, Saurav Das, public figures portrayed respectfully for pipeline testing only).

Living-document policy. After every phase gate we (1) append empirical results to this `.tex`, (2) rebuild the PDF, and (3) include the report update in the same phase merge to `master`.

2 Routing and Licensing Constraints

- **Text / coding LLMs:** OmniRoute when `USE_OMNIROUTE=1` + `OMNIROUTE_API_KEY`.
- **All images:** direct `.env` keys only (`FAL_API_KEY/FAL_KEY`, `GOOGLE_API_KEY`) — never OmniRoute.
- Local FLUX [dev] / Kontext [dev] remain R&D-only under BFL non-commercial terms; production-like finals use commercial fal/Gemini APIs.

3 Architecture (stable through Phase 3)

`src/comicengine/`: config, pricing, tracked clients, styles, analytics, stories, gallery, tasks. `scripts/`: phase CLIs. `dashboard/`: FastAPI + static UI. `data/usage.db`: per-call ledger. `data/style_lock.json`: bake-off + human-judge flags. `outputs/`: episode JSON, images, reports.

Evaluation stack. Pillow heuristics (existence, min dims, mean brightness/variance anti-blank) write `image_quality` into call meta-data. Analytics derives model recommendations (`data/model_recommendations.json`) — advisory only; scripts do not auto-switch.

4 Cumulative Ledger Snapshot

Source: live SQLite summary at report build time.

Unit costs that worked in practice. fal FLUX. `schne1l`: $\approx$ \$0.003/image, $\sim$ 1.5–2.4s. Gemini `gemini-3.1-flash-image-preview`: $\approx$ \$0.039/image, $\sim$ 9–12s. fal FLUX. `pro/kontext`: $\approx$ \$0.040/image, $\sim$ 11–14s. Text script generation (Claude Haiku): $\approx$ \$0.046 for the CJP episode.

Table 1: Spend and success by phase totals (ledger).

Phase	Calls	Cost (USD)	Notes
phase0	13	0.003	provider pings
phase0.5	2	0.046	CJP script
phase1	1	0.003	fal schnell proof
phase2	13	0.093	style bake-off + judge
phase3	23	0.387	refs + consistency + judge
phase4	11	0.655	Caesar + Hitler scripts
phase5	66	1.772	refs + 44 panels
total	129	2.960	incl. retries

Table 2: Top providers by estimated spend (post Phase 3).

Provider	Calls	Cost (USD)
google	10	0.273
fal	28	0.158
omniroute:anthropic	4	0.053
anthropic (direct)	4	0.046
omniroute:openai	3	0.002

5 Phase Log — What Ran and What Worked

5.1 Phase 0 — Environments & pings

Goal. Authenticate providers; prove routing. **What worked.** Anthropic Claude Haiku, OpenAI gpt-4.1-mini, OmniRoute **auto/cheap** for Anthropic and OpenAI, Gemini text via **direct** Google. **What failed / adapted.** Older Gemini Flash IDs error under “new user” model restrictions; replaced with **gemini-3.1-flash-lite** (text) and later **gemini-3.1-flash-image-preview** (image). One logged text error remains in the ledger (Gemini model mismatch during early ping). **Gate.** Success when calls land in SQLite with status ok.

5.2 Phase 0.5 — Script engine

Goal. Structured bedtime episode JSON for Stories UI. **What worked.** Claude Haiku → Pydantic episode `outputs/phase0.5/episode_cjp_origin.json`: title *The Cockroach Janta Party: When Students Ask for Fair Exams*, **12 panels**, fact sheet + **art_prompts**. Wikimedia Commons location refs fetched into `outputs/phase0.5/refs/` with attribution manifest (place/atmosphere refs, not official CJP press kits). **Cost.** ≈\$0.046 text (2 Anthropic calls in phase tag). **Evaluation.** Human-readable in dashboard Stories; suitable as Phase 1–3 prompt source.

5.3 Phase 1 — Single commercial image path

Goal. Prove fal FLUX **schnell** with tracked image category + QA. **What worked.**

Table 3: LLM auto-judge rubric totals (higher better; max domain score 10 each).

Style	Imm.	Bed.	Civic	Cost	Total
painterly_storybook	7	10	8	9	34
graphic_novel	8	7	10	9	34
watercolor_editorial	7	8	9	9	33
korean_manhwa	9	6	7	9	31

`scripts/phase1_generate.py -backend fal` → `outputs/phase1/flux_schnell_fal.png`; QA **pass**; \$0.003. **Evaluation.** Establishes the cheap default image backend for bake-offs.

5.4 Phase 2 — Style bake-off (complete; human lock pending)

Setup. Four presets × two fixed beats (Abhi-jeet+students; dad-daughter bedtime) via fal **schnell**: **korean_manhwa**, **painterly_storybook**, **graphic_novel**, **watercolor_editorial**. Eight bake-off images under `outputs/phase2/bakeoff/`; all QA scores **1.0 / pass**. Total image spend for 8 equal-cost styles: **\$0.024**.

Auto-judge conclusion. Winner: **painterly_storybook** (bedtime safety). Closest civic alternative: **graphic_novel**. Cheap good-enough: **watercolor_editorial**. Manhwa won immersion (9) but scored lowest bedtime safety (6).

Human policy (current). `human_judge_pending=true`; `preferred_while_judging=korean_manhwa` so active rendering stays webtoon-like while the operator compares finals in the Images tab. Override anytime with `STYLE_ID` or by editing `data/style_lock.json`.

Phase 2 also generated. Style grid previews and Gemini comparison sample (`style_gemini_01...`); Gemini image is ≈13× fal **schnell** unit cost.

5.5 Phase 3 — Character consistency (complete)

Goal. Expanded cast reference sheets + method bake-off for identity lock across poses. **Artifacts.** `outputs/phase3/consistency_manifest.json`, `decision.json`; code: `scripts/phase3_consistency_bakeoff.py`, `characters_phase3.py`, `client` helpers `fal_kontext_edit` / `gemini_image_with_refs`. Merged to master as commit 5078ead.

What worked (22/22, 0 failures).

- **8 ref sheets** (fal **schnell**): Abhijeet, Saurav Das, Modi, Amit Shah, Sonam Wangchuk, police, dad, daughter.

Table 4: Phase 3 method judge scores (higher better).

Method	Face	Outfit	Style	Rel.	Cost	Total
gemini_ref	8	8	8	9	7	40
flux_kontext	9	9	8	7	3	36
text_only_fal	3	3	7	9	10	32

- **14 consistency scenes** under interim style **korean_manhwa**: Abhijeet $\times 2$ scenes $\times 3$ methods; Dad/Modi $\times 2$ scenes $\times 2$ methods (**text_only_fal**, **gemini_ref**).
- All Pillow QA verdicts **pass**; $\approx \$0.387$ phase spend (images direct fal/Google; OmniRoute used only for text judge).

Decision (ship). Winner: **gemini_ref** (Gemini Flash Image + reference sheet) as default identity lock for hero cast. Reserve **flux_kontext** for hard Abhijeet closeup/speak retries when face drift appears. Use **text_only_fal** only for cheap drafts/extras (weak cross-scene identity despite perfect blank-frame QA).

Evaluation caveats (from decision notes). Local QA checks blanks/dims/brightness—not identity; Kontext under-sampled ($n=2$, Abhijeet-only); Gemini aspect ratios sometimes vary (e.g. 848×1264). Human eye review on Images tab remains recommended before episode mass generation.

5.6 Phase 4 — Multi-topic scripts (complete for gate)

Goal. Generalize Topic $\rightarrow$ episode JSON beyond CJP; ship two hard-history episodes written *through character lenses* with modern-affairs analogies (teen/young-adult didactic, user-selected framing).

What worked.

- Character bibles: **characters_caesar.py**, **characters_hitler.py** (ideology/thinking in notes).
- Generalized **generate_episode(story_key)** + registry **STORIES** (**cjp**, **et_tu_brutus**, **hitler_warning**).
- CLI **scripts/phase4_scripts.py**; OmniRoute default model **auto** (**auto/cheap** often skipped **tool_use**).
- Robustness: retries, JSON-only fallback, normalize panels (speaker lists $\rightarrow$ dialogue strings; verdict aliases).

Stories now in dashboard (beside CJP).

- **episode_cjp_origin** (phase0.5, 12 panels) — fair exams / youth protest.
- **episode_et_tu_brutus** (phase4, 16) — Caesar, Brutus, Cassius, Antony, People; analogies: fragile republic, personality cult, assassination is not a liberty reset.

Table 5: Phase 5 full-batch results (all QA pass).

Story	Panels ok	Errors	Primary method
episode_cjp_origin	12/12	0	gemini_ref
episode_et_tu_brutus	16/16	0	gemini_ref
episode_hitler_warning	16/16	0	gemini_ref $\times 11$ + kontext
total	44/44	0	Phase spend $\approx \$$

- **episode_hitler_warning** (phase4, 16) — Hitler/Goebbels/Eichmann ideology *exposed then condemned*; Jewish teen + resistance + bystander; analogies: propaganda, minority hate, bureaucracy of atrocity, legal demolition of democracy.

Cost. Phase 4 script spend $\approx \$0.655$ (11 text calls incl. retries). Cumulative ledger $\approx \$1.19$.

Ethics evaluation. Hitler episode is explicitly anti-fascist education (no gore; no Nazi glamour). Character POV is used so propaganda patterns are recognizable; Dad/victims close the frame. Requires human fact-check before any public publish.

5.7 Phase 5 — Panel batch generator (complete)

Goal. Episode JSON $\rightarrow$ panel PNGs with Phase 3 consistency defaults + retries + cost ledger.

What worked.

- **comicengine.panel_batch** + **scripts/phase5_panel_batch.py**
- Ref sheets under **outputs/phase5/refs/** (reuse Phase 3 when present; else **fal schnell**)
- Default render: **gemini_ref** (up to 3 cast refs) $\rightarrow$ QA-fail/error falls back to **flux_kontext** then **text_only_fal**
- Writes **image_path/image_method/image_qa** into episode JSON; Stories UI shows **/media/...** images

Evaluation. Perfect pipeline success rate for this batch. Kontext retries concentrated on Hitler-era panels (likely safety/policy or identity hardness) — validates the Phase 3 ship rule: **gemini** default, Kontext for hard retries. Human visual review still required for identity drift and historical care (especially atrocity education frames).

5.8 Phase 6 — Speech bubble compositor + story reader (complete)

Goal. Readable on-image dialogue/captions and a webpage that feels like reading a comic (not only a dump of PNGs).

Table 6: Phase 7 assembly outputs.

Story	Panels assembled	Artifacts
episode_cjp_origin	12	webtoon + PDF
episode_et_tu_brutus	16	webtoon + PDF
episode_hitler_warning	16	webtoon + PDF

What worked.

- Pillow compositor (`comicengine.compositor`): rounded speech bubbles (speaker labels + wrapped lines), caption bar, local Arial font.
- CLI `scripts/phase6_compose.py` wrote `outputs/phase6/<story>/panel_XX.png` and `composed_image_path` on each panel.
- 44/44 panels composed, 0 errors (offline, no API cost).
- Stories UI: cast list, fact sheet intro, “Opening/Meanwhile” bridge narration between panels, composed images via `/media/...`, dialogue + caption text beneath each panel, closing fact checks.

5.9 Phase 7 — Episode assembler (complete)

Goal. One readable artifact per episode: long vertical webtoon + printable PDF.

What worked.

- `comicengine.assembler` + `scripts/phase7_assemble.py`
- Cover strip + composed Phase 6 panels stacked at width 900 with gutter; multi-page PDF via Pillow.
- Outputs: `outputs/phase7/<story>/webtoon.png`, `episode.pdf`, `manifest.json`
- Stories pages expose webtoon preview + PDF/webtoon buttons.

5.10 Phase 7.5 — Only Image + Story Library (complete)

Goal. Store UI stories centrally; also ship bubble-free visual editions.

What worked.

- `edition='image_only'` assembler path $\rightarrow$ `outputs/phase7.5/<story>/webtoon_image_only.png` + `episode_image_only.pdf`
- CLI `scripts/phase7_5_image_only.py` for all three stories
- `comicengine.library` catalog at `outputs/library/catalog.json`
- Dashboard `/library` with Reader vs Only-Image editions; nav link sitewide

5.11 Phase 8 — Curation CLI + SQLite (complete)

Goal. Human gate: approve / reject / regenerate panels and episodes with durable status.

What worked.

- SQLite table `curation_item` in `data/usage.db` via `comicengine.curation`
- CLI `scripts/phase8_curate.py` (`seed|list|summary|approve|reject|regenerate`)
- Seeded 47 rows (3 episodes + 44 panels)
- Library UI chips + episode approve/reject; Alt-click regenerates panel (re-render + recompose)
- APIs: `/api/curation`, `/api/curation/seed`, `/api/curation/{id}`, regenerate endpoint

5.12 Phase 8.5 — Panel rating + prompt editor (complete)

Goal. Every reject/regen opens an editable art prompt; panels carry a human rating and free-text suggestions so quality is visible on the site and refreshes after regen.

What worked.

- Extended `curation_item` with rating (1–5) and suggestions
- `panel_editor_payload` + `regenerate_panel(..., prompt=, rating=, suggestions=)`; phase tagged `phase8.5`; edited `art_prompt` persisted into episode JSON
- `render_panel(..., prompt_override=)` uses the user-edited scene text
- Library + Story reader modal: stars, suggestions, prompt textarea; reject/regen refresh images via cache-busting `?v=`
- CLI `scripts/phase8_5_panel_editor.py`; APIs for panel editor GET and JSON regenerate body

5.13 Phase 8.6 — Reader auth + dual Vercel hosting (complete)

Goal. Separable reader vs admin surfaces; reviewers rate panels/stories without seeing engine stats; ship a shareable beta without putting image regen on Vercel.

What shipped.

- Local FastAPI: open owner console (`/admin`, `/stats`); public comics review at `/review` with thumbnail cards, Read&rate / PDF / Just read, and Mom Test after ratings; Reader CRM at `/reviewers`

- Vercel **comicengine-reader** (<https://comicengine-reader.vercel.app>): username/password (script + HttpOnly ce_session cookie + Upstash Redis)
- Library cards: evaluated 3:2 thumbnails, title, three actions — **Read & rate / PDF / Just read**
- Thumbnail picks (`scripts/select_thumbnails.py`): CJP panel 5, Et tu panel 11, Hitler-warning panel 1 (visual override away from spectacle beats)
- Vercel **comicengine-admin** (<https://comicengine-admin.vercel.app>): password gate; snapshot spend/ROI/tasks + **live** Upstash people/logins/feedback + canvas charts (20s poll)
- Branding: **ComicEngine** + **Beta** pill; footer © 2026 Avinash Nandiyala (reader, admin, local dashboard)
- Hosting runbook docs/HOSTING.md; shared Upstash resource on both projects
- After panel/story ratings: 14-question Mom Test questionnaire (Upstash); answers shown per person in admin live console

Ledger at gate (snapshot). Engine usage snapshot still \$2.96 / 138 calls (dominant spend remains Phase 5 image batch). Hosted reader/admin traffic is Upstash + static Vercel — not counted as OmniRoute/image ledger cost.

5.14 Phase 9 — ROI dashboard (complete)

Goal. Make spend actionable: unit economics + charts from the SQLite ledger.

What worked.

- `comicengine.roi.roi_dashboard()` aggregates category/purpose/daily spend + \$/story and \$/panel
- Dashboard page `/roi` + `/api/roi`; nav link sitewide
- Snapshot JSON `outputs/phase9/roi_snapshot.json` via `scripts/phase9_roi.py`

5.15 Phase 10 — Publishing pipeline (complete)

Goal. Export approved Library editions for Cloudflare Pages + R2; keep live auth/feedback on FastAPI.

What worked.

- `comicengine.publish.export_beta_site()` → `outputs/phase10/site`
- CLI `scripts/phase10_publish.py` (`export|upload-r2|deploy-pages|gate`)
- Live Google login remains on the FastAPI host (Tunnel / public HTTPS); Pages is the static mirror

5.16 Phase 11 — Cost guardrails (complete)

Goal. Freeze beta spend rules before Version 2; defer expensive cheaper-backend bake-offs until Mom Test priorities are clear.

What worked.

- Snapshot `outputs/phase11/cost_guardrails.json` with ROI baseline + reviewer/admin regen policy
- Explicit roadmap: **Version 2 begins at Phase 12**, feature work driven by exported Mom Test feedback

5.17 Romance pipeline — How I Met Your Mother Ep1 Infatuation (Steps 1–5 complete)

Goal. Dad→daughter love-story episode (korean_manhwa) with Maya teasing cuts and protagonist micro-feelings (butterflies, electricity, hand-hover, Limbo Date, feather).

Steps (TaskObserver stepper).

1. `phase_romance_step1` — script locked (**76 panels**) at `outputs/phase4/episode_how_i_met_your_mother_ep1.json`
2. `phase_romance_step2` — manhwa character refs (8/8 cast sheets)
3. `phase_romance_step3` — panel batch (`STYLE_ID=korean_manhwa`, mostly `gemini_ref`): **76/76 ok**
4. `phase_romance_step4` — speech bubbles / captions composed (76/76)
5. `phase_romance_step5` — webtoon + PDF + Library (`outputs/phase7/.../webtoon.png`, image-only under `phase7.5`)

Dashboard. Home Romance pipeline stepper + live ledger line (total spend + romance-phase spend) polled every 2s with KPIs.

Story. Father tells daughter how he met her mother (Rohan / Elena / Meridian Soft Bay trip).

Cost gate. Pre-romance ledger ≈\$2.96 / 138 calls. After Steps 1–5: romance pipeline spend ≈\$2.88 (refs \$0.018 + panels \$2.86; compose/assemble local \$0); full ledger ≈\$5.84 / 223 calls.

6 Version 2B (parallel 3D-previs track)

Version 2B is a **separate program** from Version 2 (phases 12–20) and Version 2A (sketch-first). Do not

edit `data/v2_program.json`, `data/v2a_program.json`, `data/v2a/`, or `episode_schema.py` from 2B work. Fixture is V1 HIMYM Episode 1 (76 panels); 2B owns `data/v2b/episodes/ep01.json`. Image traffic is local ComfyUI or direct `.env` APIs — never OmniRoute.

B0 architecture freeze. Docs +
`data/v2b_program.json` + dashboard /v2b. No
 Blender/ComfyUI code in that freeze.

B1 vertical slice (2026-08-27). HIMYM ep1 panel 1 living-room two-shot: headless Blender 5.2.1 Cycles (seed 42, 32 samples, 768×1152) → ComfyUI HTTP `img2img` (SD 1.5 `v1-5-pruned-emaonly`, denoise 0.42, euler 18 steps, seed 42) on Mac MPS. Local API spend \$0 (no Gemini/fal). Two Cycles beauty reruns: pixel MAE 0.0 (PNG byte hashes differ; encoder metadata). Stylizer still restages (speech-bubble/cloak drift) — expected without ControlNet; B2 pins geometry.

B2 ControlNet lock (2026-08-27). Same panel, three cameras (`cam_a` hero two-shot, `cam_b` closer, `cam_c` slight profile). Blender 5.2.1 Freestyle-as-pass is empty; lineart is Grease Pencil Line Art on black. Depth is inverted normalized Z (near=white). Index RGB: dad=R, maya=G. ComfyUI SD 1.5 multi-ControlNet on MPS (depth strength 0.75, lineart 0.65, denoise 0.45, seed 42). Local API spend \$0. SDXL waits until ControlNet fits.

Cheap skimage structure (no Depth-Anything): mean SSIM(depth) 0.575 vs hypothesis 0.7 — calibrated floor 0.53. Mean edge IoU 0.368. Visual geometry lock holds (sofa + two figures, no B1 restage).

Camera	SSIM(depth)	edge IoU
<code>cam_a</code>	0.541	0.416
<code>cam_b</code>	0.513	0.293
<code>cam_c</code>	0.671	0.397

B3 style LoRA lock (2026-08-27). Same three cameras, no Blender rerender. Acquired `neonforestmist/sd15-storybook-anime-lora` (CreativeML OpenRAIL-M; SHA256 `0304bb42...e9b0`) because SDXL+ControlNet does not fit 16GB MPS. Locked on every ControlNet stylize: LoRA strength 1.0, checkpoint `v1-5-pruned-emaonly`, euler 18, CFG 6.5, seed 42, denoise 0.65 (raised from B2 0.45 so the adapter shows through). Registry: `data/v2b/lora/registry.json`. Weights gitignored. Structure holds: mean SSIM(depth) 0.584 (calibrated floor 0.53); mean edge IoU 0.292; mean pixel MAE vs B2 0.047. Local API spend \$0. Gut-check is these three framings, not a 10-scene VLM (B8). Style reads as storybook illustration on capsule meshes (warmer lamp, white spherical heads); B4 replaces capsules with block humanoids + Dad LoRA.

Camera	SSIM(depth)	edge IoU	MAE vs B2
<code>cam_a</code>	0.559	0.276	0.039
<code>cam_b</code>	0.525	0.286	0.054
<code>cam_c</code>	0.669	0.314	0.048

G1 eval harness (2026-08-27). Prove-shot eval on the same three cameras. No PPO/GRPO/DPO. Structure hard gate is **mean** SSIM(depth) ≥ 0.53 (not per-camera; `cam_b` is often 0.51–0.52). Locked B3 mean SSIM 0.584 — pass. Identity vs beauty is cheap `grid_hist_8x8` (log-only on G1 capsules; B4 uses DINOv2). Pairwise judge: `gemini-3.6-flash` direct `GOOGLE_API_KEY` (retired `gemini-2.5-flash` returned 404). A/B swap; never OmniRoute images. 12/12 human labels on `/v2b/gate1`. Exact LLM–human 4/12; tie-lenient 10/12. Two hard disagreements: `b_b2_beauty` and `c_b2_beauty` (human prefers beauty, Gemini prefers B2). Exact 8/12 is calibration debt, not a train signal. Harness still passes. Best-of-N seeds {42, 43, 44, 45} on existing AOVs. Winners: `cam_a` seed 44, `cam_b` seed 44, `cam_c` seed 42 (copy of locked B3). Human and Gemini both pick locked B3 over the `cam_b` BoN winner. 46 successful Gemini pairwise calls, \$0.043. Preferences: `data/v2b/eval/preferences.jsonl`. Scorecard: `data/v2b/eval/himym_ep01_g1_scorecard.json`. B8 stays locked.

Pair	Gemini	Human
<code>a_b1_b2</code>	B	B
<code>a_b2_b3</code>	B	tie
<code>a_b1_b3</code>	B	B
<code>a_b3_beauty</code>	A	tie
<code>b_b2_b3</code>	B	tie
<code>b_b2_beauty</code>	A	B
<code>b_b3_beauty</code>	A	tie
<code>b_bon_b3</code>	B	B
<code>c_b2_b3</code>	B	tie
<code>c_b2_beauty</code>	A	B
<code>c_b3_beauty</code>	A	tie
<code>c_bon_b3</code>	tie	tie

B4 character LoRA (2026-08-27). Capsules cannot be a character dataset. B4 replaces them with block humanoids (Dad: navy sweater + curly hair spheres; Maya: hoodie + ponytail) and writes a sibling tree `outputs/v2b/himym_ep01/b4/` so G1 `cam_{a,b,c}` stays frozen. Standing Dad turntable (no sofa), `-quick` 8 views $\times$ 2 seeds. Bootstrap stylize: denoise 0.40, ControlNet depth 0.55 / lineart 0.45, style LoRA on, trigger `ce_dad_rohan`. Four Maya standing views as contrast only — no Maya LoRA (that is B6). Train: SD 1.5 rank-16 LoRA in `ComfyUI/.venv` (diffusers+PEFT, 250 steps on 12 cached latents). Compass named `kohya/SDXL`; this 16GB Mac cannot. Weights gitignored; SHA256 `bfd7f7d8...3bfaf` in `data/v2b/lora/registry.json`. Infer: stacked `LoraLoader` (style then Dad at 0.8) on the new living-room AOVs. Local API spend \$0. Never OmniRoute. Scorecard `data/v2b/eval/himym_ep01_b4.json`: mean

SSIM(depth) 0.5488 (floor 0.53) — pass. Restage MAE vs beauty 0.061–0.082 (≤ 0.12). DINOv2-small on 4 hold-out stylized views: 3/4 closer to Dad sheet than Maya — pass (–quick bar; full bar is 6/8). `grid_hist_8x8` stays log-only; do not claim Compass 0.85. `cam_b` SSIM 0.469 fails per-camera (same pattern as G1); the hard gate is the mean. Gut-check: hair mass readable on `cam_a`; Maya is a smaller hoodie figure, not an orange capsule. B8 stays locked.

Camera	SSIM(depth)	edge IoU	MAE vs beauty
<code>cam_a</code>	0.534	0.263	0.061
<code>cam_b</code>	0.469	0.244	0.082
<code>cam_c</code>	0.644	0.252	0.063

B5 location reuse (2026-08-27). Panel specs are versioned JSON under `data/v2b/specs/` (Pydantic; not `episode_schema.py`). Headless `build_scene.py` places primitives from that JSON. Prove: living room $\times 3$ cameras (stacked style+Dad LoRA) and Grand Oriole lobby $\times 2$ (style only, empty — Young Rohan has no LoRA). Sibling tree `outputs/v2b/himym_ep01/b5/`. G1/B4 stay frozen. Local API spend \$0. Never OmniRoute. Location eval paints object-index characters out, then DINOv2-small on backgrounds: living intra-mean 0.825, lobby pair 0.737, living–lobby 0.476. Same-room beats cross-room — pass. Compass 0.9 is a hypothesis, not claimed. Structure floor 0.53 is the living-room calibration: living mean SSIM(depth) 0.5489 — pass (`cam_b` 0.469 still under per-camera). Lobby mean 0.478 (wide 0.388) — empty-room cheap-depth, logged separately. Restage MAE vs beauty ≤ 0.082 living / ≤ 0.039 lobby. B8 stays locked. Tram waits.

Panel	SSIM(depth)	MAE vs beauty	location
<code>living_a</code>	0.534	0.061	<code>living_room</code>
<code>living_b</code>	0.469	0.082	<code>living_room</code>
<code>living_c</code>	0.644	0.063	<code>living_room</code>
<code>lobby_wide</code>	0.388	0.039	<code>grand_oriole_lobby</code>
<code>lobby_close</code>	0.569	0.038	<code>grand_oriole_lobby</code>

B6 multi-character masks (2026-08-27). Maya gets the same SD 1.5 turntable LoRA recipe as Dad (`ce_maya`, rank 16, 250 steps, SHA256 `b767153f...8d86`). Prove: living-room `cam_a` and `cam_c` only — `cam_b` is the hard crop. Two-pass: style+Dad on the full frame, then stock ComfyUI `VAEEncodeForInpaint` on the G object-index with Maya LoRA (no global Dad+Maya stack). Sibling tree `outputs/v2b/himym_ep01/b6/`. G1/B4/B5 stay frozen. Local API spend \$0. Never OmniRoute. No InstantID. Maya holdout DINOv2-small 4/4 closer to Maya sheet than Dad — pass (–quick bar). Structure mean SSIM(depth) 0.5917 (floor 0.53) — pass. Restage MAE vs beauty 0.065–0.067. Bleed: Dad R-crop vs pass1 0.966 / 0.984 (floor 0.90) — the mask held. Maya G-crop own > other both cameras. Seated `dad_own` vs Maya sheet is a coin-flip on

block meshes in a tight two-shot — log-only, same honesty as SSIM 0.53 vs Compass 0.7. Compass 0.85 not claimed. B8 stays locked.

Camera	SSIM(depth)	MAE vs beauty	Dad vs pass1	Maya own>
<code>cam_a</code>	0.537	0.065	0.966	yes
<code>cam_c</code>	0.647	0.067	0.984	yes

7 Dashboard Surfaces

Local (<http://127.0.0.1:8765>): usage KPIs, TaskObserver, Stories, Library, ROI, Analytics, Images, admin console, public /review, Report PDF (/report).

Hosted beta: reader <https://comicengine-reader.vercel.app>; fellowdev admin <https://comicengine-admin.vercel.app> (password-gated). Live image regen / OmniRoute scripting remain on the laptop (or a GPU/API host).

8 What We Explicitly Learned

1. OmniRoute is excellent for LLM scripting/judging; keeping images off OmniRoute avoided opaque image billing and matched commercial-key policy.
2. fal `schnell` makes wide style exploration financially trivial (~\$0.024 for 8 bake-off frames).
3. Auto style judges trade off immersion vs. bedtime safety; human adjudication remains mandatory for product voice. G1 exact Gemini–human agreement was 4/12 (tie-lenient 10/12) — humans tied often; Gemini preferred stylized B3 over beauty and over B2.
4. Pillow “pass” is necessary but not sufficient for character consistency; the Phase 3 LLM judge correctly ranked `gemini_ref` over cheaper text-only despite identical QA scores.
5. OmniRoute model routing matters: prefer `auto` (or provider-prefixed Claude) when `auto/cheap` returns thinking text without `tool_use`.
6. Vercel serverless session UX needs `HttpOnly` cookies (Bearer-only + aggressive /me clears caused login loops); `CSS display on .btn` overrides HTML `[hidden]`.
7. Version 2B B5: DINOv2 background cosine separates living-room cameras (0.825) from Grand Oriole lobby (0.476 cross); Compass 0.9 is still a hypothesis.
8. Version 2B B6: two-pass G-index inpaint keeps Dad vs pass1 at 0.966+; Maya holdout 4/4. Seated `dad_own` vs Maya sheet on block two-shots is log-only.

9 Risks and Ethics

CJP bedtime lane remains ages ~7–10; Phase 4 hard-history lane is teen didactic (~13–17). Hitler content must

never romanticize Nazism; Caesar content must not romanticize political murder. Respectful depiction of public figures; Commons attribution for place photos; human fact-check before any public publish. Style and consistency locks must not be decided solely by LLM judges or blank-frame QA. Public beta URLs expose comics widely — keep admin password private; Upstash holds reader hashes/sessions/feedback.

10 Next Gates

- Human style lock: clear `human_judge_pending` (manhwa vs painterly).
- Collect panel/story ratings from beta readers; harvest admin live people analytics for Version 2 priorities.
- Re-run `prepare_admin_site.py` (+ redeploy) whenever ledger/ROI snapshot must refresh.
- Version 2B B7: storyboard sequencing with cache. B8 stays locked.

11 Conclusion

ComicEngine now covers
 script→render→compose→assemble→curate→ROI→hosted
 reader/admin beta→publish packaging. Version 2 starts
 at Phase 12 using logged panel/story feedback from beta
 readers.

A Reproduce

```
source .venv/bin/activate
pip install -e .
# configure .env (never commit)
python scripts/run_dashboard.py
python scripts/build_report.py
# PDF: http://127.0.0.1:8765/report
```

B Changelog (report editions)

2026-08-09 a

Initial two-column architecture draft.

2026-08-09 b

Living phase log: ledger tables, Phase 2 judge matrix, Phase 3 interim notes.

2026-08-09 c

Phase 3 gate closed: 22/22 images, winner `gemini_ref` (judge 40), ledger \$0.532; dashboard `/report` link.

2026-08-09 d

Phase 4 scripts: Et tu Brutus + Hitler warning (teen didactic, character POV) beside CJP; ledger ≈\$1.19.

2026-08-09 e

Phase 5: 44/44 panels rendered (`gemini_ref/Kontext`); Stories show images; ledger ≈\$2.96.

2026-08-09 f

Phase 6: speech bubbles/captions on 44 panels; reader bridges + cast/facts on Stories pages.

2026-08-09 g

Phase 7: vertical webtoon + PDF for all three stories; Stories download/preview links.

2026-08-09 h

Phase 7.5: Only Image editions + Story Library (`/library`) catalog.

2026-08-09 i

Phase 8: curation SQLite + CLI + Library approve/reject/regenerate.

2026-08-09 j

Phase 9: ROI dashboard (`/roi`) unit economics + spend charts.

2026-08-09 k

Phase 8.5: panel rating + editable art prompt on reject/regen.

2026-08-09 l

Phase 8.6: dual Vercel reader/admin + Upstash auth; thumbnails; live people analytics; Beta branding; ledger snapshot still ≈\$2.96 / 138 calls.

2026-08-27 a

Version 2B B0/B1: 3D-previs architecture freeze; HIMYM ep1 panel 1 Cycles→ComfyUI SD 1.5; beauty seed MAE 0.0.

2026-08-27 b

Version 2B B2: three-camera AOVs + SD 1.5 ControlNet; mean SSIM(depth) 0.575 (calibrated floor 0.53); visual geometry lock.

2026-08-27 c

Version 2B B3: locked SD 1.5 storybook style LoRA (strength 1.0, denoise 0.65); mean SSIM(depth) 0.584; mean MAE vs B2 0.047.

2026-08-27 d

Version 2B G1: 12/12 human pairs; Gemini exact 4/12 (tie-lenient 10/12); BoN seeds 42–45; \$0.043 `gemini-3.6-flash`; no DPO.

2026-08-27 e

Version 2B B4: block humanoids; Dad SD 1.5 LoRA (250 steps, SHA256 `bfd7f7d8...`); mean SSIM 0.549; DINOv2 3/4 Dad>Maya; B8 still locked.

2026-08-27 f

Version 2B B5: spec-driven living room ×3 + Grand Oriole lobby ×2; living SSIM 0.549; DINOv2 same-room 0.825 > cross 0.476; no Rohan LoRA.

2026-08-27 g

Version 2B B6: Maya SD 1.5 LoRA (250 steps) + two-pass
G-index inpaint on cam_a/c; SSIM 0.592; holdout 4/4;
Dad vs pass1 0.966+; B8 still locked.